



Università degli Studi di Messina

Data Science



Big Data Acquisition Spark

Prof. Daniele Ravi



dravi@unime.it

Recap from Last Lecture

- MapReduce → new distributed computing framework suitable for working with large scale datasets
- Useful in all those situations where data need to be accessed **sequentially**
- May be hard to program and does not support well multiple map-reduce rounds

Data-Flow Systems

- MapReduce uses 2 "ranks" of tasks: one for **map** the other for **reduce**
- Data flows from the first rank to the second
- Generalized Data-Flow Systems abstract from this in two ways:
 - Allow any number of "ranks"/tasks
 - Allow functions other than just map and reduce
- As long as data goes in one direction only, recovery at intermediate rank is possible

MapReduce: Criticisms

- **3 major limitations of MapReduce:**
 - Hard to program directly
 - Many problems are not easily described as map-reduce
 - Relies heavily on disk I/O communication which can become a bottleneck for performance.
 - **Latency:** MapReduce is optimized for batch processing and may not be suitable for real-time applications.
 - Persistence to disk slower than in-memory computation
 - MapReduce is not suitable for more complex operations composed of several map-reduce steps or iterative algorithms (e.g., machine learning)

Apache Spark Overview

- Apache Spark is an open-source, distributed computing system designed for big data processing.
- It provides an advanced data analytics engine that can handle large-scale data with speed and scalability.
- Spark was developed in response to **the limitations of Hadoop's MapReduce** model, focusing on making data processing faster and more efficient.
- It's particularly suited for large-scale data processing tasks like batch processing, streaming, machine learning, and graph computation.

Spark: Introduction

- Originally developed at UC Berkeley in 2009 and later donated to the Apache Software Foundation (open-source)
- Implemented in **Scala** (running on top of the Java Virtual Machine)
- Unified **computing engine** (**Spark Core**)
- Set of **high-level APIs** for data analysis:
 - **Spark SQL** (structured data), **MLlib** (machine learning), **GraphX** (graph analytics),
Spark Streaming (stream data processing)

Spark: Introduction

- Unlike Hadoop, Spark does not come with a storage system
- In fact, it provides interfaces for many local and distributed storage systems:
 - HDFS, Amazon S3, Cassandra, Hive Metastore, or classical RDBMS
- Additionally, Spark's APIs are available for many programming languages:
Scala, Java, Python, and R
- This flexibility is the key of its success in the Big Data landscape

Spark: Most Popular Data-Flow System

- Computing framework not limited to map-reduce model
- In addition to MapReduce, Spark provides:
 - Fast data sharing (no intermediate saving to local disks + caching)
 - General execution graphs (DAGs)
 - Richer functions than just map and reduce
- Compatible with Hadoop

Key Features of Spark

1. **Speed:** Spark uses in-memory computation (keeping data in RAM rather than reading and writing from disk), which makes it faster than Hadoop MapReduce for many tasks.
 - It can also spill data to disk when memory is insufficient.
2. **Ease of Use:** Spark offers APIs in several programming languages
3. **Unified Engine:** Spark integrates several libraries into one unified framework allow for ease of implementation in various domains: **Spark SQL, MLLib , GraphX, Spark Streaming**
4. **Fault Tolerance:** Like Hadoop, Spark is resilient to failures. It uses Directed Acyclic Graphs (DAGs) to represent job execution plans, and its **lineage-based fault tolerance** allows lost data to be recomputed in case of node failure.
5. **Scalability:** Spark can scale from a single machine to thousands of nodes in a cluster, handling petabytes of data.

Lineage-based fault tolerance

- Is a key concept in frameworks like **Apache Spark**.
- It refers to the ability of the system to **recompute lost data** by keeping track of the **lineage** of data transformations, instead of replicating data across nodes.
- **Here's how it works:**
 - **Lineage:** Each dataset is generated by applying a sequence of transformations to an initial dataset (e.g., map, filter, reduce).
 - This sequence of transformations is referred to as the **lineage** of the dataset.
 - Spark keeps track of this lineage of operations, so it knows how a dataset was derived from its original source.

Lineage-based fault tolerance

- **Fault Tolerance via Lineage:**
 - Rather than restarting the entire computation or replicating data in advance, Spark can **recompute** only the lost part of the data by **re-executing** the transformations based on the lineage information.
 - This is more efficient especially in cases of intermittent failure.
- **Directed Acyclic Graph (DAG):**
 - Spark represents the lineage of transformations as a **DAG**.
 - Each node in this graph represents an RDD (Resilient Distributed Dataset), and the edges represent transformations applied to those RDDs.
 - If any part of the computation fails, Spark can trace the lineage back through the DAG and recompute only the necessary parts.

How Apache Spark Works

- **RDD (Resilient Distributed Dataset)**
 - It is a distributed collection of objects across the nodes in a cluster, allowing for parallel processing.
 - RDDs are **immutable**, meaning once created, they cannot be altered.
 - However, transformations can be applied to them to create new RDDs.
 - They are fault-tolerant, and Spark can reconstruct lost data by tracking the transformations that led to each RDD.
- **Transformations:** Operations like `map()`, `filter()`, `flatmap()` are lazily evaluated and do **not execute immediately**. They generate a new RDD and build up a DAG.
- **Actions:** Operations like `reduce()`, `collect()`, and `save()` trigger the execution of the transformations and bring results back to the driver program or write them to external storage.

How Apache Spark Works

- **Directed Acyclic Graph (DAG):**
 - When a Spark job is submitted, Spark constructs a DAG to represent the series of transformations applied to the data.
 - It optimizes the execution plan before running the job, minimizing data shuffling and communication between nodes.
- **In-Memory Computation:**
 - Spark performs operations in-memory, avoiding the overhead of disk I/O seen in traditional MapReduce systems.
 - This leads to much faster execution times, especially for iterative tasks like machine learning algorithms or data analytics.

How Apache Spark Works

- **Job Execution:**
 - A **driver** program coordinates the tasks and transformations.
 - **Workers** execute the tasks in parallel on their respective partitions of the data.
 - Tasks are distributed across nodes in a cluster, and intermediate data can be cached in memory for faster access.
- **Cluster Manager:**
 - Spark can work with different cluster managers for resource allocation:
 - **Standalone mode:** Spark's built-in manager.
 - **Apache YARN:** Used in Hadoop ecosystems.
 - **Mesos:** General-purpose cluster manager.
 - **Kubernetes:** Container orchestration platform.

Summary: Spark Features

- Fault-tolerant system
- In-memory caching which enables efficient execution of multi-round algorithms (i.e., multiple sequential tasks)
 - performance improvement w.r.t. Hadoop
- Spark can run:
 - on a single machine → local mode
 - on a cluster managed by a cluster manager (e.g., Spark Standalone, YARN, Mesos)

Spark: Features

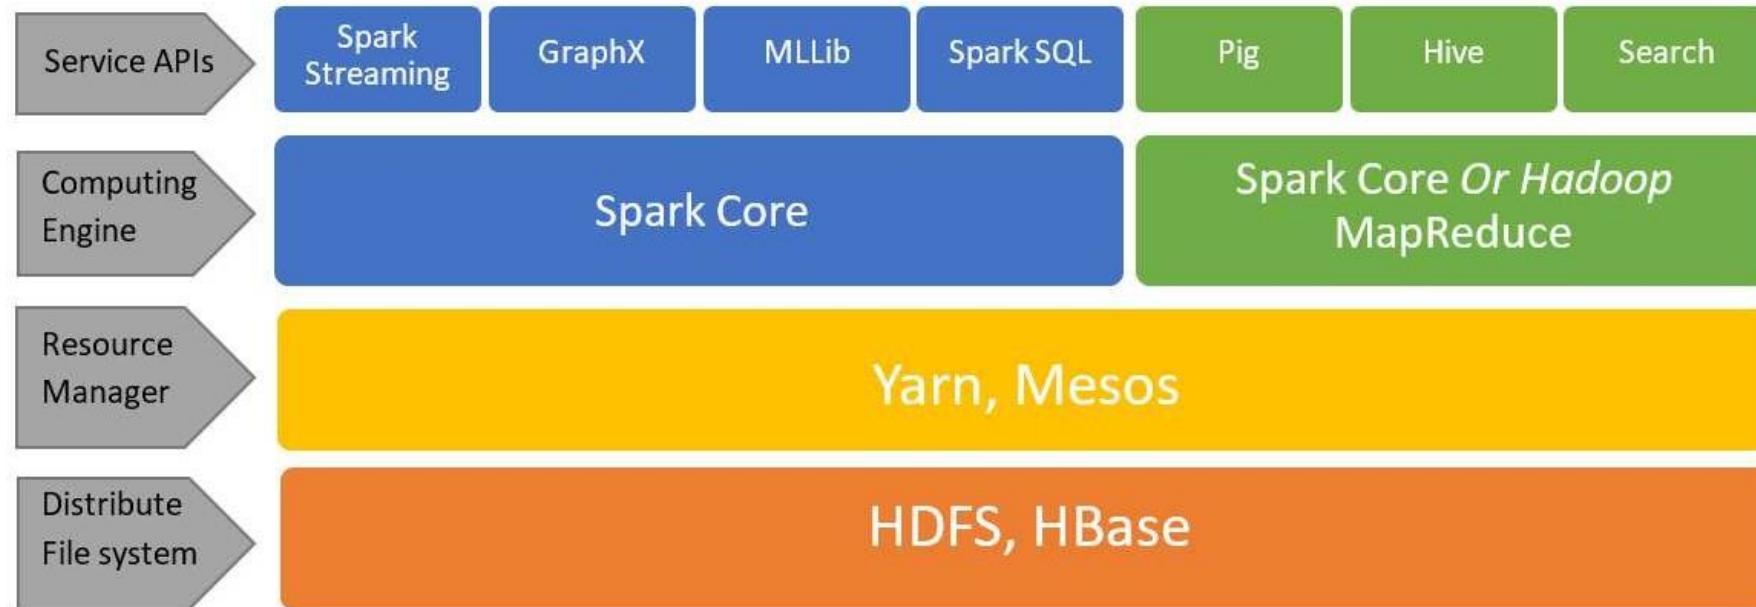


Figure 1 - Spark Context

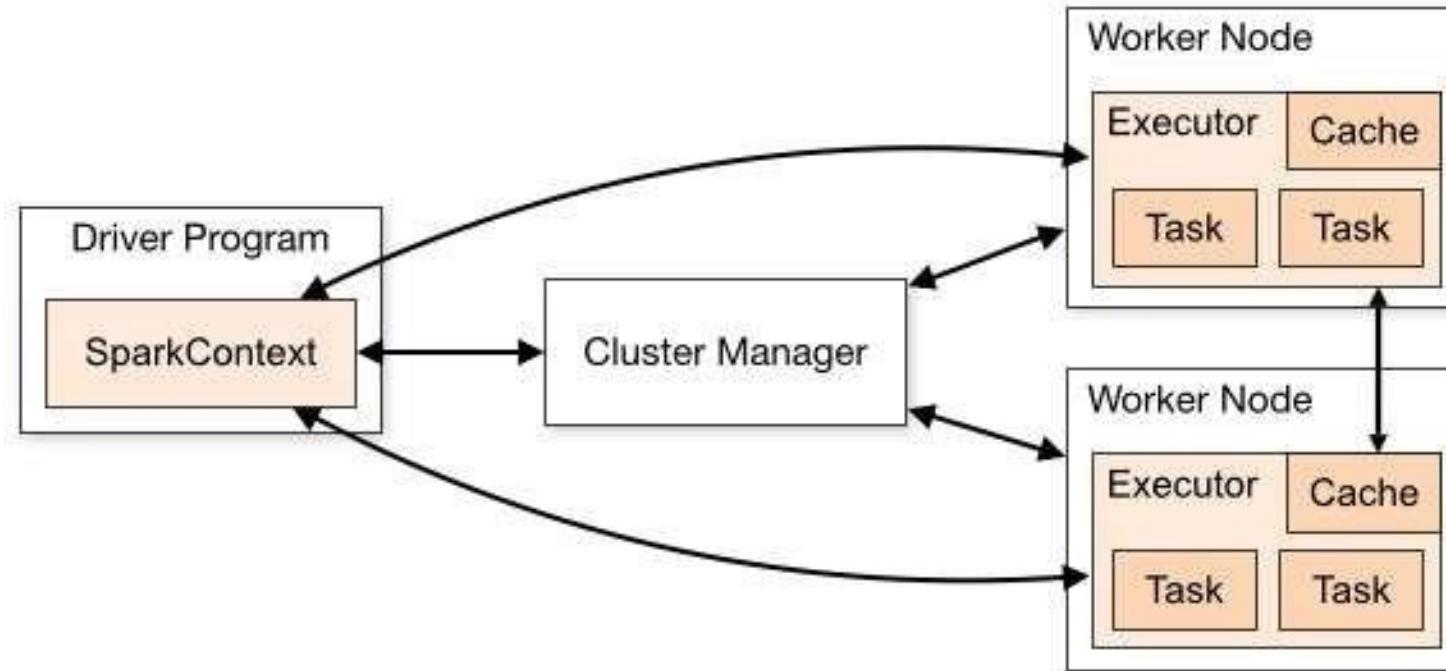
Spark Application: Driver

- The **driver process** (a.k.a. **master** in MapReduce terminology) runs the application's entry point from a node in the cluster
- The driver is responsible for:
 - Maintaining information about the application
 - Responding to a user program or input
 - Analyzing, distributing, and scheduling work across executors
- The driver is represented by an object called **Spark Context**

Spark Application: Executor(s) and Cluster Manager

- **Executor** processes (a.k.a. **workers** in Hadoop terminology) actually compute the tasks assigned by the driver
- Each executor is responsible for:
 - Running the code assigned to it by the driver
 - Reporting the state of the computation back to the driver
- The **cluster manager** controls physical machines and allocates resources to applications

Spark Application



Spark Application: Considerations

- Driver and executors are processes which can live on the same machines or on different nodes
- When Spark is running in local mode, both the driver and executors are running as separate threads on the same machine
- Executors mostly run Scala code
- Driver can be governed by different languages using Spark's APIs

Resilient Distributed Dataset (RDD)

- Fundamental **abstraction** of Spark to indicate a collection of elements of the same type
 - Generalization of MapReduce's key-value pairs
- RDDs are **partitioned** and possibly spread across multiple nodes of the cluster
- Best suited for applications that apply the same operation across all the elements of the dataset

RDD: Partitions

- Each RDD is split into chunks called partitions distributed across nodes
- A program can specify the number of partitions for an RDD (otherwise Spark will choose one)
- Programmer can also decide whether to use the default Hash Partitioner or a custom one
- A typical number of partitions is 2 or 3 times the number of **total cores in the cluster**

RDD: Partitions

- Partitioning enables the following:
 - **Data reuse - >** data is kept in executors' main memory so as to avoid expensive access to external disks
 - **Parallelism - >** Some data transformations are applied independently to each partition thereby avoiding expensive data transfers

RDD: Characteristics

- RDDs are **immutable** (i.e., read-only)
- Can be created either from data stored on distributed file system (e.g., HDFS) or as a result of transformations of other RDDs
- RDDs do not need to be always materialized
 - Each RDD maintains a sort of "trace" of transformations (lineage) that led to the current status
 - This way, RDD can always be re-created even upon a failure

RDD Operations

- Let A be an RDD, the following **3 operations** are possible:
 - **Transformations** - > generate a new RDD B from the data in A
 - **Actions** - > launch a computation on the data in A, which returns a value to the application
 - **Persistence** - > save the RDD in memory for later actions

RDD Operations: Transformations

- **Narrow:** each partition of A contributes at most to one partition of B (e.g., `map`)
 - No need to transfer data across nodes
- **Wide:** each partition of A may contribute to multiple partitions of B (e.g., `groupBy`)
 - Possible need to transfer data across nodes
- **Lazy evaluation:** nothing is computed unless required by an action

Narrow Transformations

In Apache Spark, transformations are classified as either **narrow** or **wide** depending on how data moves between partitions during execution.

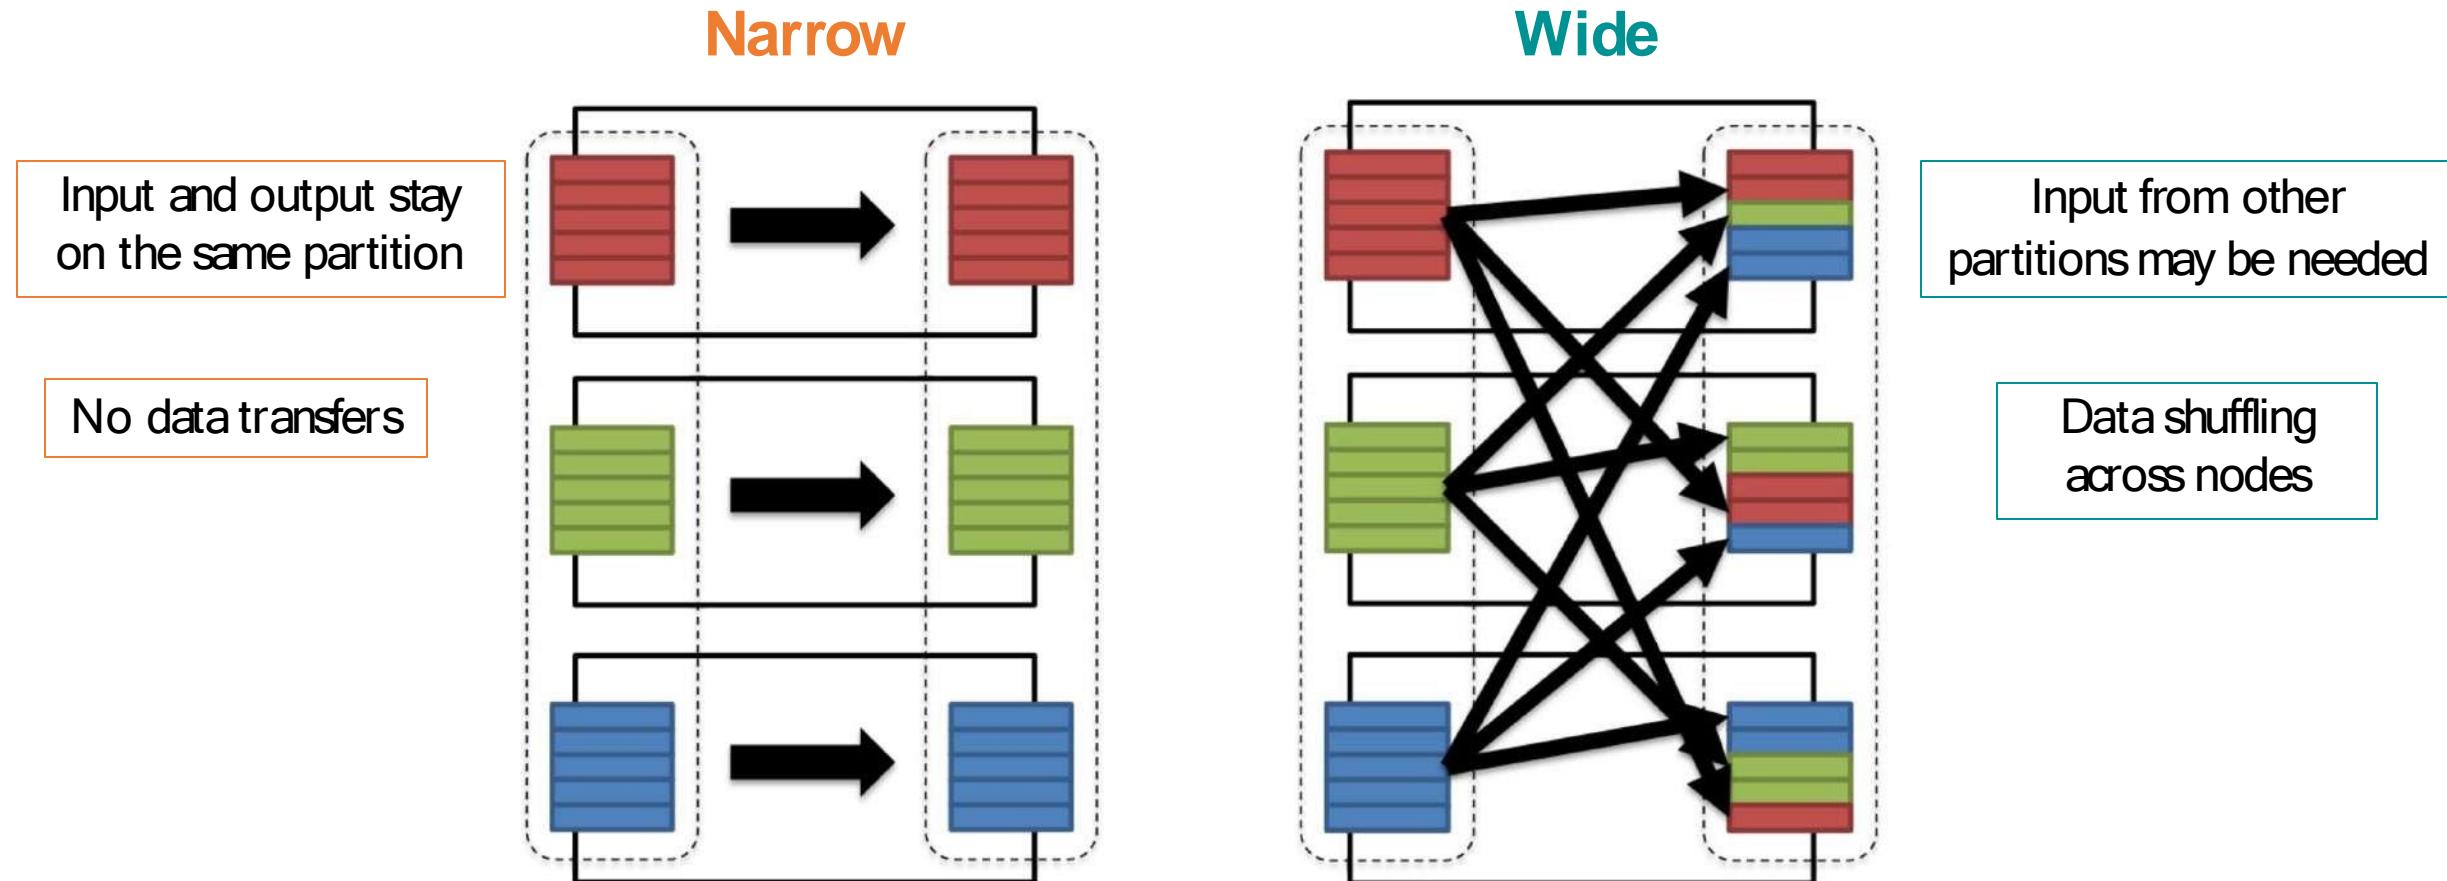
Narrow Transformations

- Are operations where each partition of the input RDD is used by at most one partition of the output RDD.
- **Characteristics:**
 - Data dependencies are localized within partitions, which makes processing efficient.
 - Examples include transformations like `map()`, `filter()`, and `mapPartitions()`.
- **Data Flow:** There is no need for data to be transferred across the cluster, so narrow transformations are generally faster.

Wide Transformations

- Are operations where each partition of the input RDD may be used by multiple partitions of the output
- **Characteristics:**
 - These transformations require data to be reshuffled across partitions, which involves network communication, making them more time-consuming.
 - Examples include transformations like `reduceByKey()`, `groupByKey()`, and `join()`.
- **Data Flow:** Data needs to be moved across the nodes in the cluster, which involves combining and reorganizing data.
- **Use Cases:** Situations where data from different partitions must be brought together, for example, to perform aggregations or joins.
- **Summary**
 - Involve data shuffling and network I/O, making them more resource-intensive.
 - Efficient use of transformations in Spark involves minimizing wide transformations, as they typically represent a major bottleneck due to the need for shuffling data across the cluster.

Narrow vs. Wide Transformations



RDD Operations: Actions

- **Example:** the count method returns the number of elements of the RDD
- When the action is called the RDD is actually materialized (lazy evaluation)

Spark DataFrame and Dataset APIs

- RDDs are the most basic data model used by Spark
 - low-level and schema-less
- On top of RDD API, **Spark SQL** module provides **2 interfaces** to operate on structured data like tables in relational databases:
 - **DataFrame API**
 - **Dataset API**

Spark: DataFrame

- Distributed collection of data organized into **named columns**
- Allows higher level abstraction than plain vanilla RDDs
- Similar to [Pandas DataFrame](#) unless few differences
- Dataset API is available only for Scala and Java as it extends DataFrame API with type-safe, object-oriented programming interface

Spark DataFrame vs. Pandas DataFrame

- Spark DataFrames are **immutable**: once created they cannot be modified
- As for RDDs, Spark may apply 2 kinds of operations on DataFrames:
transformations and **actions**
- Lazy evaluation allows to queue transformations applied to elements of a DataFrame until an action is called
- DataFrame (and Dataset aswell) can be turned back to RDD

Spark vs. Hadoop MapReduce

- **Performance:** Spark is usually faster
 - In-memory data processing vs. data persistencing to disk after any map/reduce step
 - Spark requires lots of memory to run fast, otherwise its performance deteriorates
 - MapReduce integrates better with other services running
- **Ease of use:** Spark provides a higher-level API which is easier to program
- **Data processing:** Spark is more flexible and general

Useful Extra References

- Spark official <https://spark.apache.org/>
- What is Spark? [<https://www.databricks.com/spark/about>]
- Getting Started with Spark [<https://www.databricks.com/spark/getting-started-with-apache-spark>]
- Spark tutorial [<https://data-flair.training/blogs/spark-tutorial/>]
- PySpark Tutorial (video)
[<https://www.youtube.com/watch?v= C8kWso4ne4>]
- Docker con spark [https://hub.docker.com/_/spark]

Take-Home Message of Today

- Spark is a general-purpose distributed data processing engine which overcomes many of the Hadoop's limitations
- Spark provides a rich ecosystem of services to work on (big) data through APIs accessible via multiple programming languages
- Lazy execution avoids frequent and costly disk operations
- Spark's **DataFrame** as the main abstraction for playing with data

CODING with SPARK

Overview of some essential code in Spark

1. Setting Up a `SparkSession`

``SparkSession`` is the entry point for data processing.

```
from pyspark.sql import SparkSession
```

```
# Create a SparkSession
```

```
spark = SparkSession.builder.appName("BasicSparkExample").getOrCreate()
```

- ``appName("BasicSparkExample")``: Sets the name for the application.
- ``getOrCreate()``: Either creates a new ``SparkSession`` or returns an existing one.

Overview of some essential code in Spark

2. Creating an RDD (Resilient Distributed Dataset)

RDDs are the core data structure in Spark.

```
# Create an RDD from a Python list  
data = [1, 2, 3, 4, 5]  
rdd = spark.sparkContext.parallelize(data)
```

```
# Show the contents of the RDD  
print(rdd.collect())
```

- `spark.sparkContext.parallelize(data)`: Converts a Python list to an RDD.
- `rdd.collect()`: Collects all the elements of the RDD to the driver program.
- **Note:** For large datasets, avoid `collect()` as it might cause memory issues.

Overview of some essential code in Spark

- **3. Performing Transformations on an RDD**

Transformations create a new RDD from an existing one.

```
# Map transformation - multiply each element by 2  
mapped_rdd = rdd.map(lambda x: x * 2)
```

```
# Filter transformation - keep only even numbers  
filtered_rdd = mapped_rdd.filter(lambda x: x % 2 == 0)
```

```
# Collect the results  
print(filtered_rdd.collect()) # Output: [4, 8]
```

- `map(lambda x: x * 2)`: Applies the given function to each element.
- `filter(lambda x: x % 2 == 0)`: Filters the elements based on the given condition.

Overview of some essential code in Spark

- **4. Loading Data Using `SparkSession`**
- Spark is often used to process large files, like CSVs or JSON.

```
# Load a CSV file
csv_path = "/path/to/your/data.csv" # Replace with the path to your CSV file
df = spark.read.option("header", True).csv(csv_path)
```

```
# Show the first 5 rows of the DataFrame
df.show(5)
```

- `.read.option("header", True).csv(csv_path)`: Reads a CSV file into a DataFrame with headers.
- `.show(5)`: Displays the first 5 rows of the DataFrame.

Overview of some essential code in Spark

- **5. Performing Transformations on DataFrames**
- Transformations can also be applied to DataFrames. Here are examples of selecting columns and filtering rows:

```
# Select specific columns  
selected_df = df.select("name", "age")
```

```
# Filter rows where age is greater than 30  
filtered_df = selected_df.filter(df.age > 30)
```

```
# Show the result  
filtered_df.show()
```

- `df.select("name", "age")`: Selects specific columns from the DataFrame.
- `df.filter(df.age > 30)`: Filters rows based on the condition.

Overview of some essential code in Spark

- **6. Aggregating Data with GroupBy**

We can group data by a column and calculate aggregate metrics:

```
# Group by the "age" column and count the number of records for each age  
grouped_df = df.groupBy("age").count()
```

```
# Show the result  
grouped_df.show()
```

- `df.groupBy("age").count()`: Groups the data by the `age` column and counts the occurrences for each unique age.

Overview of some essential code in Spark

- **7. Actions on DataFrames**

Actions trigger the actual computation in Spark.

```
# Count the number of rows in the DataFrame  
row_count = df.count()  
print(f"Number of rows: {row_count}")
```

```
# Collect data to the driver  
collected_data = df.collect()  
print(collected_data)
```

- `.`count()`: Returns the number of rows in the DataFrame.
- `.`collect()`: Collects the DataFrame rows to the driver as a list.
- Note: Be cautious when using `.`collect()` for large datasets.

Overview of some essential code in Spark

- **8. Writing Data to Output**

You can save your DataFrame to various formats such as CSV :

```
# Save the filtered DataFrame as a CSV file  
output_path = "/path/to/output/folder"  
filtered_df.write.option("header", True).csv(output_path)
```

- `df.write.option("header", True).csv(output_path)` : Saves the DataFrame as a CSV file to the specified path.

Summary

- `SparkSession` is the entry point for working with Spark.
- `RDDs`: Low-level data structure that you can create using `'sparkContext'` and apply transformations like `'map'` and `'filter'`.
- `DataFrames`: Higher-level abstraction for data, similar to tables. You can load data from files, perform column operations, group by values, and write results.
- Transformations (like `'map()'`, `'filter()'`, `'groupBy()'`) are lazy operations used to manipulate data.
- Actions (like `'count()'`, `'collect()'`) trigger actual computation.

Common Transformations in Spark

- In Apache Spark, transformations are operations that create a new dataset from an existing one.
- They are lazy, meaning that they don't immediately compute their results.
- Instead, transformations are recorded and executed only when an action (like a count or collect) is called, which triggers the computation.

Common Transformations in Spark

1. map(func)

- Applies a function to each element of the RDD/DataFrame and returns a new RDD/DataFrame.
 - **Example:** Converting a list of numbers to their squares.

```
rdd = sc.parallelize([1, 2, 3, 4])
squared = rdd.map(lambda x: x ** 2)
# Result: [1, 4, 9, 16]
```

2. flatMap(func)

- Similar to `map()`, but the output is flattened. Each input item can be mapped to 0 or more items, and the results are combined into a single RDD.
 - **Example:** Splitting lines of text into words.

```
rdd = sc.parallelize(["Hello world", "Spark is great"])
words = rdd.flatMap(lambda line: line.split(" "))
# Result: ["Hello", "world", "Spark", "is", "great"]
```

Common Transformations in Spark

3. filter(func)

- Returns a new RDD/DataFrame containing only the elements that satisfy the given condition
- **Example:** Filtering even numbers.

```
rdd = sc.parallelize([1, 2, 3, 4, 5, 6])
even = rdd.filter(lambda x: x % 2 == 0)
# Result: [2, 4, 6]
```

4. distinct()

- Returns a new RDD/DataFrame with duplicate elements removed.
- **Example:** Removing duplicates from a list.

```
rdd = sc.parallelize([1, 2, 2, 3, 4, 4, 5])
unique = rdd.distinct()
# Result: [1, 2, 3, 4, 5]
```

Common Transformations in Spark

5. union(otherDataset)

Returns a new RDD/DataFrame that contains the union of the elements from the two datasets.

- **Example:** Combining two lists.

```
rdd1 = sc.parallelize([1, 2, 3])
rdd2 = sc.parallelize([4, 5, 6])
combined = rdd1.union(rdd2)
# Result: [1, 2, 3, 4, 5, 6]
```

6. intersection(otherDataset)

- Returns a new RDD/DataFrame containing only the elements present in both datasets (intersection).
- **Example:** Finding common elements between two lists.

```
rdd1 = sc.parallelize([1, 2, 3, 4])
rdd2 = sc.parallelize([3, 4, 5, 6])
common = rdd1.intersection(rdd2)
# Result: [3, 4]
```

Common Transformations in Spark

7. groupByKey()

- Groups all values associated with each key in an RDD of key-value pairs.
- **Example:** Grouping numbers by key.

```
rdd = sc.parallelize([("a", 1), ("b", 2), ("a", 3)])
grouped = rdd.groupByKey()
# Result: [("a", [1, 3]), ("b", [2])]
```

8. reduceByKey(func)

- Combines values with the same key using the specified function.
- **Example:** Adding values for the same key.

```
rdd = sc.parallelize([("a", 1), ("b", 2), ("a", 3)])
summed = rdd.reduceByKey(lambda x, y: x + y)
# Result: [("a", 4), ("b", 2)]
```

Common Transformations in Spark

9. **join(otherDataset)**

- Joins two RDDs/DataFrames by their keys.
- **Example:** Joining two datasets by common key.

```
rdd1 = sc.parallelize([("a", 1), ("b", 2)])
rdd2 = sc.parallelize([("a", 3), ("b", 4)])
joined = rdd1.join(rdd2)
    # Result: [("a", (1, 3)), ("b", (2, 4))]
```

10. **cartesian(otherDataset)**

- Returns the Cartesian product of two datasets.
- **Example:** Generating all possible pairs.

```
rdd1 = sc.parallelize([1, 2])
rdd2 = sc.parallelize([3, 4])
product = rdd1.cartesian(rdd2)
    # Result: [(1, 3), (1, 4), (2, 3), (2, 4)]
```

Common Transformations in Spark

11. **sortBy(func, ascending=True)**

- Sorts the RDD by the given function.
- **Example:** Sorting numbers in ascending order.

```
rdd = sc.parallelize([4, 2, 7, 1])
sorted_rdd = rdd.sortBy(lambda x: x)
# Result: [1, 2, 4, 7]
...  
...
```

12. **sample(withReplacement, fraction)**

- ****Description**:** Returns a sampled subset of the dataset.
- **Example:** Sampling elements from an RDD.

```
rdd = sc.parallelize(range(10))
sampled_rdd = rdd.sample(False, 0.4)
# Sample about 40% of the elements
```

Scenario 1

- We have a dataset containing customer transactions (`transactions.csv`), and we want to:
 1. Load the dataset.
 2. Filter transactions to only include purchases over \$100.
 3. Group by customer ID and sum their total purchase amounts.
 4. Output the result.
- **Sample Data (`transactions.csv`):**
- Each log entry is a space-separated string with the following format:
`transaction_id,customer_id,amount,timestamp`

1,101,150.75,2023-09-20 10:15:00

2,102,80.50,2023-09-20 11:20:00

3,101,200.00,2023-09-21 09:45:00

4,103,55.00,2023-09-21 12:10:00

5,104,500.30,2023-09-22 15:25:00

Generate Random Data on a CSV file

```
import csv
import random
from datetime import datetime, timedelta
import uuid

# Numero di transazioni da generare
num_transactions = 1000

# Lista per memorizzare le transazioni
transactions = []

# Data di inizio per le timestamp (ultimi 30
giorni)
end_date = datetime.now()
start_date = end_date - timedelta(days=30)

# Generazione delle transazioni
for _ in range(num_transactions):
    transaction = {
        'transaction_id': str(uuid.uuid4()), # ID univoco
per ogni transazione
        'customer_id': random.randint(1000, 9999), # ID cliente casuale
        'amount': round(random.uniform(10.0, 1000.0),
2), # Importo casuale tra 10 e 1000
        'timestamp': (start_date + timedelta(
            seconds=random.randint(0, int((end_date -
start_date).total_seconds()))))
            ).strftime('%Y-%m-%d %H:%M:%S')
    }
    transactions.append(transaction)
```

Generate Random Data on a CSV file

```
# Ordinamento delle transazioni per timestamp
transactions.sort(key=lambda x: x['timestamp'])

# Scrittura su file CSV
filename = 'transactions.csv'
with open(filename, 'w', newline='') as csvfile:
    fieldnames = ['transaction_id', 'customer_id', 'amount', 'timestamp']
    writer = csv.DictWriter(csvfile, fieldnames=fieldnames)

    writer.writeheader()
    for transaction in transactions:
        writer.writerow(transaction)

print(f"File {filename} generato con successo con {num_transactions} transazioni!")
```

Implementation

```
from pyspark.sql import SparkSession
from pyspark.sql import functions as F

# Initialize Spark session
spark = SparkSession.builder.appName("CustomerTransactionsAnalysis").getOrCreate()

# Step 1: Load the dataset
df = spark.read.csv("transactions.csv", header=True, inferSchema=True)

# Step 2: Filter transactions where the amount is greater than 100
filtered_df = df.filter(df['amount'] > 100)
```

Implementation

```
# Step 3: Group by customer ID and sum the total amount spent
grouped_df = filtered_df.groupBy("customer_id").agg(F.sum("amount").alias("total_spent"))
    # The .agg() function is used to perform an aggregate operation on the grouped data.
    # Here, F.sum("amount") calculates the total of the "amount" column for each group of customer_id.
    # F is an alias for the pyspark.sql.functions module, which provides a set of functions for manipulating and
    # transforming data.
    # alias("total_spent") is used to rename the resulting column to "total_spent".
```



```
# Step 4: Show the final result
grouped_df.show()
```



```
# (Optional) Save the result to a CSV file
grouped_df.write.csv("customer_totals.csv", header=True)
```



```
# Stop the Spark session
spark.stop()
```

Scenario 2 (Machine Learning MLlib)

- You have a dataset of customer information (e.g., `customer_data.csv`) and you want to predict whether a customer will buy a product (binary outcome: 1 for purchase, 0 for no purchase) based on features like age, income, and gender
- **Sample `customer_data.csv`:**
 - Each log entry is a comma-separated string with the following format:
`age,income,gender,label`

25,50000,0,0

32,60000,1,1

47,120000,0,1

51,130000,1,0

30,55000,0,0

Implementation

```
from pyspark.sql import SparkSession
from pyspark.ml.feature import VectorAssembler
from pyspark.ml.classification import LogisticRegression
from pyspark.ml.evaluation import BinaryClassificationEvaluator

# Step 1: Initialize Spark session
spark = SparkSession.builder.appName("LogisticRegressionExample").getOrCreate()

# Step 2: Load the data
data = spark.read.csv("customer_data.csv", header=True, inferSchema=True)

# Step 3: Assemble features into a single vector
assembler = VectorAssembler(inputCols=["age", "income", "gender"], outputCol="features")
data_with_features = assembler.transform(data)
final_data = data_with_features.select("features", "label")
```

Implementation

```
# Step 4: Split the data into training and test sets
```

```
train_data, test_data = final_data.randomSplit([0.8, 0.2])
```

```
# Step 5: Train a Logistic Regression model
```

```
logistic_regression = LogisticRegression(featuresCol='features', labelCol='label')
```

```
lr_model = logistic_regression.fit(train_data)
```

```
# Step 6: Make predictions on the test set
```

```
test_results = lr_model.transform(test_data)
```

```
# Step 7: Evaluate the model performance
```

```
evaluator = BinaryClassificationEvaluator(labelCol="label", rawPredictionCol="rawPrediction")
```

```
accuracy = evaluator.evaluate(test_results)
```

```
print(f"Test Accuracy = {accuracy}")
```

```
# (Optional) Save the model
```

```
lr_model.save("logistic_regression_model")
```

```
# Stop the Spark session
```

```
spark.stop()
```

Scenario 3: Word Count with Spark Streaming

- To create a word count program using Apache Spark Streaming in Python, you'll first need to set up a real-time data stream and then count words as they appear.
- In this example we make use of the `socketTextStream()` method, which listens to data on a socket.
- We can send text data to that socket through a terminal using `nc -lk <port>`
- In Spark Streaming, the `SparkContext` is typically used instead of `SparkSession`.

Scenario 3: Word Count with Spark Streaming

```
from pyspark import SparkContext
from pyspark.streaming import StreamingContext

# Create a local StreamingContext with two working threads and batch interval of 1 second
sc = SparkContext("local[2]", "NetworkWordCount")
ssc = StreamingContext(sc, 1)

# Connect to hostname:port, e.g., localhost:9999
lines = ssc.socketTextStream("localhost", 9999)

# Split each line into words
words = lines.flatMap(lambda line: line.split(" "))

# Count each word in each batch
pairs = words.map(lambda word: (word, 1))
word_counts = pairs.reduceByKey(lambda x, y: x + y)

# Print the word counts
word_counts.pprint()

# Start computation and wait for the streaming to finish
ssc.start()
ssc.awaitTermination()
```

Scenario 3: Word Count with Spark Streaming

Step-by-Step Explanation:

1. Imports:

1. Import the necessary Spark components (SparkContext, StreamingContext).

2. Setup Contexts:

1. Create a SparkContext and a StreamingContext.
2. The StreamingContext has a batch interval of 1 second, meaning it processes data in 1-second chunks.

3. Create the Stream:

1. `ssc.socketTextStream("localhost", 9999)` sets up the stream to receive data on port 9999.

4. Process the Data:

1. `flatMap(lambda line: line.split(" "))` splits each line into individual words.
2. `map(lambda word: (word, 1))` pairs each word with the value 1.
3. `reduceByKey(lambda x, y: x + y)` counts the occurrences of each word.

5. Output:

1. `pprint()` prints the word counts.

6. Start Streaming:

1. `ssc.start()` starts the computation.
2. `ssc.awaitTermination()` keeps the program running until you manually stop it.

Scenario 3: Word Count with Spark Streaming

Running the Application

Start a socket listener using Netcat:

```
nc -l k 9999
```

This command will open a terminal where you can type messages that Spark will process.

Run the Spark script using:

```
spark-submit word_count.py
```

Scenario 4: Spark SQL

This is a simple example of using Spark SQL in Python to create a DataFrame, register it as a temporary table, and then query it using SQL.

```
from pyspark.sql import SparkSession  
# Initialize SparkSession  
spark = SparkSession.builder .appName("Spark SQL Example") .getOrCreate()  
  
# Create a DataFrame with sample data  
data = [  
    ("Alice", 34),  
    ("Bob", 45),  
    ("Cathy", 29),  
    ("David", 40)  
]
```

Scenario 4: Spark SQL

```
columns = ["Name", "Age"]
```

```
df = spark.createDataFrame(data, columns)
```

```
# Register the DataFrame as a temporary SQL table  
df.createOrReplaceTempView("people")
```

```
# Use Spark SQL to select people with age greater than 30  
result = spark.sql("SELECT Name, Age FROM people WHERE Age > 30")
```

```
# Show the result  
result.show()
```

```
# Stop the Spark session  
spark.stop()
```

Scenario 4: Spark SQL

Explanation

- 1. Initialize SparkSession:** We create a SparkSession which allows us to interact with Spark.
- 2. Create DataFrame:** A DataFrame is created with sample data (data and columns).
- 3. Register Temporary Table:** The DataFrame is registered as a temporary SQL table using createOrReplaceTempView.
- 4. Query with Spark SQL:** We run an SQL query using the spark.sql method to select rows where Age is greater than 30.
- 5. Show the result:** The result.show() function prints the output.