

# Files Formats – not just CSV

- Key factor in Big Data processing and query performance
- Schema Evolution
- Compression and Splittability
- Data Processing
  - Write performance
  - Partial read
  - Full read

# Optimizations – CPU and I/O

Statistics for filtering and query optimization

projection push down

X	Y	Z
x1	y1	z1
x2	y2	z2
x3	y3	z3
x4	y4	z4
x5	y5	z5


predicate push down

X	Y	Z
x1	y1	z1
x2	y2	z2
x3	y3	z3
x4	y4	z4
x5	y5	z5


read only the data you need

X	Y	Z
x1	y1	z1
x2	y2	z2
x3	y3	z3
x4	y4	z4
x5	y5	z5

Minimizes CPU cache misses




cache misses costs cpu cycles



# Parquet

- columnar storage format
- key strength is to store nested data in truly columnar format using definition and repetition levels<sup>1</sup>



(1) Dremel made simple with parquet - <https://blog.twitter.com/2013/dremel-made-simple-with-parquet>

# What is a Parquet File?

Apache Parquet is a columnar storage file format designed for efficient data processing and analytics.

- **Key Features:**

- **Columnar Format:** Stores data by columns, not rows.
- **Optimized for Big Data:** Ideal for distributed systems like Apache Hadoop, Spark, and Hive.
- **Compression:** Supports efficient compression and encoding schemes to reduce file size.

- **Advantages:**

- 1. **Efficient I/O:** Faster queries by reading only required columns.
- 2. **Small File Size:** Highly compressed, saving storage.
- 3. **Schema Evolution:** Supports changes to schema over time.
- 4. **Interoperability:** Works seamlessly with data tools like Hive, Spark, and Pandas.



# Key Features of Parquet

- **Columnar Storage:**

- Ideal for analytical workloads where specific columns are queried frequently.

- **Data Compression:**

- Reduces storage using advanced encoding like Snappy, GZIP, and LZO.

- **Metadata Support:**

- Stores schema and statistics for efficient querying.

- **Language-Independent:**

- Can be used with Java, Python, C++, and more.

- **Parquet vs Other Formats**

Feature	Parquet	CSV	JSON	Avro
File Size	Small	Large	Medium	Small
Read Efficiency	High	Low	Medium	Medium
Schema Support	Yes	No	No	Yes
Compression	Advanced	None	None	Basic

# Common Use Cases

## 1. Big Data Analytics:

- Storing and processing large datasets with tools like Hadoop and Spark.

## 2. Data Warehousing:

- Optimized for OLAP workloads.

## 3. Cloud Storage:

- Supported by AWS S3, Google Cloud Storage, and Azure.

## 4. ETL Pipelines:

- Efficient for extracting, transforming, and loading data.

## • Read/Write Parquet Files:

- **Python (Pandas):** `import pandas as pd df = pd.read_parquet('file.parquet') # Reading  
df.to_parquet('output.parquet') # Writing`

- **Apache Spark:** `spark.read.parquet("file.parquet") # Reading df.write.parquet("output.parquet") # Writing`


## • Storage Tools:

- Hive, Presto, and Snowflake support native Parquet formats.



# Parquet file structure & Configuration

Internal structure of parquet file



Configurable parquet parameters

Property name	Default value	Description
parquet.block.size	128 MB	The size in bytes of a block (row group).
parquet.page.size	1MB	The size in bytes of a page.
parquet.dictionary.page.size	1MB	The maximum allowed size in bytes of a dictionary before falling back to plain encoding for a page.
parquet.enable.dictionary	true	Whether to use dictionary encoding.
parquet.compression	UNCOMPRESSED	The type of compression: UNCOMPRESSED, SNAPPY, GZIP & LZO

Parquet is state-of-the-art, open-source columnar format that supports *most* of Hadoop processing frameworks and is optimized for high compression and high scan efficiency

# Encoding

- Delta Encoding:
  - E.g timestamp can be encoded by storing first value and the delta between subsequent values which tend to be small due to temporal validity
- Prefix Encoding:
  - delta encoding for strings
- Dictionary Encoding:
  - Small set of values, e.g post code, ip addresses etc
- Run Length Encoding:
  - repeating data

# Delta Encoding:

- **Definition:** Delta encoding reduces storage requirements by encoding data as the difference (or "delta") between sequential values instead of storing the actual values.
- **Example:**
  - Given timestamps: [1000, 1005, 1010, 1015]
  - Instead of storing all timestamps, delta encoding stores:
    - First value: 1000
    - Differences: [+5, +5, +5]
  - This reduces redundancy because differences are typically smaller in magnitude, allowing for more efficient compression.
- **Use Case:** Commonly used in time-series data, where values are sequential and changes between consecutive data points are small.

# Prefix Encoding:

- **Definition:** A variation of delta encoding, applied specifically to strings. It identifies common prefixes and replaces them with a shorter representation.
- **Example:**
  - Strings: ["apple", "applet", "applesauce"]
  - Encoding:
    - Base: "apple"
    - Additions: ["", "t", "sauce"]
  - This encodes only the differences from the base string.
- **Use Case:** Effective in applications with similar string data, such as filenames, URLs, or text with repetitive phrases.

# Dictionary Encoding:

- **Definition:** Compresses data by replacing repeated values with a reference to an entry in a dictionary (a predefined mapping of values to unique identifiers).
- **Example:**
  - Data: ["CA", "NY", "CA", "CA", "TX"]
  - Dictionary:
    - "CA" -> 1
    - "NY" -> 2
    - "TX" -> 3
  - Encoded Data: [1, 2, 1, 1, 3]
- **Use Case:** Useful for categorical data with a limited set of unique values, such as postal codes, country codes, or IP addresses.

# Run-Length Encoding (RLE):

- **Definition:** Compresses data by replacing consecutive repeated values with a single value and a count of repetitions.
- **Example:**
  - Data: [A, A, A, B, B, C, C, C, C]
  - Encoded: [(A, 3), (B, 2), (C, 4)]
- **Use Case:** Effective for datasets with long runs of the same value, such as image data with large regions of the same color, or simple binary data with repeating patterns.

# Comparison and Applications:

- **Delta Encoding** is ideal for ordered numeric data with small incremental changes, like sensor readings.
- **Prefix Encoding** is best for string data with common prefixes.
- **Dictionary Encoding** works well for datasets with repetitive discrete values.
- **Run-Length Encoding** excels with repetitive sequences in data, especially in multimedia compression.

# AVRO

- Language neutral data serialization system
  - Write a file in python and read it in C
- AVRO data is described using language independent schema
- AVRO schemas are usually written in JSON and data is encoded in binary format
- Supports schema evolution
  - producers and consumers at different versions of schema
- Supports compression and are splittable

# Avro – File structure and example

Sample AVRO schema in JSON format

```
{  
    "type" : "record",  
    "name" : "tweets",  
    "fields" : [ {  
        "name" : "username",  
        "type" : "string",  
    }, {  
        "name" : "tweet",  
        "type" : "string",  
    }, {  
        "name" : "timestamp",  
        "type" : "long",  
    } ],  
    "doc" : "schema for storing tweets"  
}
```

Avro file structure

